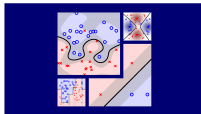


Machine Learning Techniques (機器學習技法)



Lecture 10: Random Forest

Hsuan-Tien Lin (林軒田)

htlin@csie.ntu.edu.tw

Department of Computer Science
& Information Engineering

National Taiwan University
(國立台灣大學資訊工程系)



Roadmap

- 1 Embedding Numerous Features: Kernel Models
- 2 Combining Predictive Features: Aggregation Models

Lecture 9: Decision Tree

recursive branching (purification) for **conditional aggregation** of **constant hypotheses**

Lecture 10: Random Forest

- Random Forest Algorithm
- Out-Of-Bag Estimate
- Feature Selection
- Random Forest in Action

- 3 Distilling Implicit Features: Extraction Models

Recall: Bagging and Decision Tree

Bagging

function **Bag**(\mathcal{D}, \mathcal{A})

For $t = 1, 2, \dots, T$

- ① request size- N' data $\tilde{\mathcal{D}}_t$ by **bootstrapping** with \mathcal{D}
- ② obtain base g_t by $\mathcal{A}(\tilde{\mathcal{D}}_t)$

return $G = \text{Uniform}(\{g_t\})$

—**reduces variance**

by voting/averaging

Decision Tree

function **DTree**(\mathcal{D})

if **termination** return base g_t

else

- ① learn $b(\mathbf{x})$ and split \mathcal{D} to \mathcal{D}_c by $b(\mathbf{x})$
- ② build $G_c \leftarrow \text{DTree}(\mathcal{D}_c)$
- ③ return $G(\mathbf{x}) =$

$$\sum_{c=1}^C \mathbb{I}[b(\mathbf{x}) = c] G_c(\mathbf{x})$$

—**large variance**

especially if fully-grown

putting them together?

(i.e. **aggregation of aggregation :-)**)

Random Forest (RF)

random forest (RF) = bagging + fully-grown C&RT decision tree

function **RandomForest**(\mathcal{D})

For $t = 1, 2, \dots, T$

- ① request size- N' data $\tilde{\mathcal{D}}_t$ by **bootstrapping** with \mathcal{D}
- ② obtain tree g_t by **DTree**($\tilde{\mathcal{D}}_t$)

return $G = \text{Uniform}(\{g_t\})$

function **DTree**(\mathcal{D})

if **termination** return base g_t

else

- ① learn $b(\mathbf{x})$ and split \mathcal{D} to \mathcal{D}_c by $b(\mathbf{x})$
- ② build $G_c \leftarrow \text{DTree}(\mathcal{D}_c)$
- ③ return $G(\mathbf{x}) =$

$$\sum_{c=1}^C \mathbb{I}[b(\mathbf{x}) = c] G_c(\mathbf{x})$$

- highly **parallel/efficient** to learn
- **inherit pros** of C&RT
- **eliminate cons** of fully-grown tree

Diversifying by Feature Projection

recall: **data randomness** for **diversity** in **bagging**

randomly **sample N' examples** from \mathcal{D}

another possibility for **diversity**:

randomly **sample d' features** from \mathbf{x}

- when sampling index $i_1, i_2, \dots, i_{d'}$: $\Phi(\mathbf{x}) = (x_{i_1}, x_{i_2}, \dots, x_{i_{d'}})$
- $\mathcal{Z} \in \mathbb{R}^{d'}$: a **random subspace** of $\mathcal{X} \in \mathbb{R}^d$
- often $d' \ll d$, efficient for large d
—can be generally applied on other models
- original RF **re-sample new subspace for each $b(\mathbf{x})$ in C&RT**

RF = **bagging** + **random-subspace C&RT**

Diversifying by Feature Expansion

randomly **sample d' features** from \mathbf{x} : $\Phi(\mathbf{x}) = \mathbf{P} \cdot \mathbf{x}$
with **row i of \mathbf{P}** sampled randomly \in **natural basis**

more **powerful** features for **diversity**: row i other than natural basis

- **projection** (combination) with random row \mathbf{p}_i of \mathbf{P} : $\phi_i(\mathbf{x}) = \mathbf{p}_i^T \mathbf{x}$
- often consider **low-dimensional** projection:
only **d'' non-zero** components in \mathbf{p}_i
- includes **random subspace** as **special case**:
 $d'' = 1$ and $\mathbf{p}_i \in$ **natural basis**
- original RF consider d' random **low-dimensional projections for each $b(\mathbf{x})$** in C&RT

RF = **bagging** + random-**combination** C&RT
—**randomness** everywhere!

Fun Time

Within RF that contains random-combination C&RT trees, which of the following hypothesis is equivalent to each branching function $b(\mathbf{x})$ within the tree?

- ① a constant
- ② a decision stump
- ③ a perceptron
- ④ none of the other choices

Fun Time

Within RF that contains random-combination C&RT trees, which of the following hypothesis is equivalent to each branching function $b(\mathbf{x})$ within the tree?

- ① a constant
- ② a decision stump
- ③ a perceptron
- ④ none of the other choices

Reference Answer: ③

In each $b(\mathbf{x})$, the input vector \mathbf{x} is first projected by a random vector \mathbf{v} and then thresholded to make a binary decision, which is exactly what a perceptron does.

Bagging Revisited

Bagging

function **Bag**(\mathcal{D}, \mathcal{A})

For $t = 1, 2, \dots, T$

- ① request size- N' data $\tilde{\mathcal{D}}_t$
by **bootstrapping** with \mathcal{D}
- ② obtain base g_t by $\mathcal{A}(\tilde{\mathcal{D}}_t)$

return $G = \text{Uniform}(\{g_t\})$

	g_1	g_2	g_3	\dots	g_T
(\mathbf{x}_1, y_1)	$\tilde{\mathcal{D}}_1$	*	$\tilde{\mathcal{D}}_3$		$\tilde{\mathcal{D}}_T$
(\mathbf{x}_2, y_2)	*	*	$\tilde{\mathcal{D}}_3$		$\tilde{\mathcal{D}}_T$
(\mathbf{x}_3, y_3)	*	$\tilde{\mathcal{D}}_2$	*		$\tilde{\mathcal{D}}_T$
\dots					
(\mathbf{x}_N, y_N)	$\tilde{\mathcal{D}}_1$	$\tilde{\mathcal{D}}_2$	*		*

* in t -th column: not used for obtaining g_t
—called **out-of-bag (OOB) examples** of g_t

Number of OOB Examples

OOB (in \star) \iff not sampled after N' drawings

if $N' = N$

- probability for (\mathbf{x}_n, y_n) to be OOB for g_t : $(1 - \frac{1}{N})^N$
- if N large:

$$\left(1 - \frac{1}{N}\right)^N = \frac{1}{\left(\frac{N}{N-1}\right)^N} = \frac{1}{\left(1 + \frac{1}{N-1}\right)^N} \approx \frac{1}{e}$$

OOB size per $g_t \approx \frac{1}{e}N$

OOB versus Validation

OOB

	g_1	g_2	g_3	\dots	g_T
(\mathbf{x}_1, y_1)	\tilde{D}_1	★	\tilde{D}_3		\tilde{D}_T
(\mathbf{x}_2, y_2)	★	★	\tilde{D}_3		\tilde{D}_T
(\mathbf{x}_3, y_3)	★	\tilde{D}_2	★		\tilde{D}_T
\dots					
(\mathbf{x}_N, y_N)	\tilde{D}_1	★	★		★

Validation

	g_1^-	g_2^-	\dots	g_M^-
$\mathcal{D}_{\text{train}}$	$\mathcal{D}_{\text{train}}$	$\mathcal{D}_{\text{train}}$		$\mathcal{D}_{\text{train}}$
\mathcal{D}_{val}	\mathcal{D}_{val}	\mathcal{D}_{val}		\mathcal{D}_{val}
\mathcal{D}_{val}	\mathcal{D}_{val}	\mathcal{D}_{val}		\mathcal{D}_{val}
$\mathcal{D}_{\text{train}}$	$\mathcal{D}_{\text{train}}$	$\mathcal{D}_{\text{train}}$		$\mathcal{D}_{\text{train}}$

- ★ like \mathcal{D}_{val} : ‘enough’ random examples unused during training
- use ★ to validate g_t ? easy, but **rarely needed**
- use ★ to validate G ? $E_{\text{oob}}(G) = \frac{1}{N} \sum_{n=1}^N \text{err}(y_n, G_n^-(\mathbf{x}_n))$,
with G_n^- contains only trees that \mathbf{x}_n is OOB of,
such as $G_N^-(\mathbf{x}) = \text{average}(g_2, g_3, g_T)$

E_{oob} : self-validation of bagging/RF


Model Selection by OOB Error

Previously: by Best E_{val}

$$g_{m^*} = \mathcal{A}_{m^*}(\mathcal{D})$$

$$m^* = \underset{1 \leq m \leq M}{\operatorname{argmin}} E_m$$

$$E_m = E_{\text{val}}(\mathcal{A}_m(\mathcal{D}_{\text{train}}))$$



RF: by Best E_{OOB}

$$G_{m^*} = \text{RF}_{m^*}(\mathcal{D})$$

$$m^* = \underset{1 \leq m \leq M}{\operatorname{argmin}} E_m$$

$$E_m = E_{\text{OOB}}(\text{RF}_m(\mathcal{D}))$$

- use E_{OOB} for **self-validation** —of RF **parameters** such as d''
- **no re-training** needed

E_{OOB} often **accurate** in practice

Fun Time

For a data set with $N = 1126$, what is the probability that $(\mathbf{x}_{1126}, y_{1126})$ is not sampled after bootstrapping $N' = N$ samples from the data set?

- 1 0.113
- 2 0.368
- 3 0.632
- 4 0.887

Fun Time

For a data set with $N = 1126$, what is the probability that $(\mathbf{x}_{1126}, y_{1126})$ is not sampled after bootstrapping $N' = N$ samples from the data set?

- ① 0.113
- ② 0.368
- ③ 0.632
- ④ 0.887

Reference Answer: ②

The value of $(1 - \frac{1}{N})^N$ with $N = 1126$ is about 0.367716, which is close to $\frac{1}{e} = 0.367879$.

Feature Selection

for $\mathbf{x} = (x_1, x_2, \dots, x_d)$, want to remove

- **redundant** features: like keeping one of 'age' and 'full birthday'
- **irrelevant** features: like insurance type for cancer prediction

and only 'learn' **subset-transform** $\Phi(\mathbf{x}) = (x_{i_1}, x_{i_2}, x_{i_{d'}})$

with $d' < d$ for $g(\Phi(\mathbf{x}))$

advantages:

- **efficiency**: simpler hypothesis and shorter prediction time
- **generalization**: 'feature noise' removed
- **interpretability**

disadvantages:

- **computation**: 'combinatorial' optimization in training
- **overfit**: 'combinatorial' selection
- **mis-interpretability**

decision tree: a rare model
with **built-in feature selection**

Feature Selection by Importance

idea: if possible to calculate

importance(i) for $i = 1, 2, \dots, d$

then can select $i_1, i_2, \dots, i_{d'}$ of top- d' **importance**

importance by linear model

$$\text{score} = \mathbf{w}^T \mathbf{x} = \sum_{i=1}^d w_i x_i$$

- intuitive estimate: **importance**(i) = $|w_i|$ with some 'good' **w**
- getting 'good' **w**: learned from data
- non-linear models? often **much harder**

next: 'easy' feature selection in RF

Feature Importance by Permutation Test

idea: random test

—if feature i needed, ‘random’ values of $x_{n,i}$ degrades performance

- which random values?
 - uniform, Gaussian, \dots : $P(x_i)$ changed
 - bootstrap, **permutation** (of $\{x_{n,i}\}_{n=1}^N$): $P(x_i)$ approximately remained
- **permutation** test:

$$\text{importance}(i) = \text{performance}(\mathcal{D}) - \text{performance}(\mathcal{D}^{(p)})$$

with $\mathcal{D}^{(p)}$ is \mathcal{D} with $\{x_{n,i}\}$ replaced by **permuted** $\{x_{n,i}\}_{n=1}^N$

permutation test: a general statistical tool for arbitrary non-linear models like RF

Feature Importance in Original Random Forest

permutation test:

$$\text{importance}(i) = \text{performance}(\mathcal{D}) - \text{performance}(\mathcal{D}^{(p)})$$

with $\mathcal{D}^{(p)}$ is \mathcal{D} with $\{x_{n,i}\}$ replaced by **permuted** $\{x_{n,i}\}_{n=1}^N$

- **performance**($\mathcal{D}^{(p)}$): needs re-training and **validation** in general
- ‘**escaping**’ **validation**? **OOB** in RF
- original RF solution: $\text{importance}(i) = E_{\text{oob}}(G) - E_{\text{oob}}^{(p)}(G)$,
where $E_{\text{oob}}^{(p)}$ comes from replacing each request of $x_{n,i}$ by a **permuted OOB** value

RF **feature selection** via **permutation** + **OOB**:
often efficient and promising in practice

Fun Time

For RF, if the 1126-th feature within the data set is a constant 5566, what would $\text{importance}(i)$ be?

- 1 0
- 2 1
- 3 1126
- 4 5566

Fun Time

For RF, if the 1126-th feature within the data set is a constant 5566, what would $\text{importance}(i)$ be?

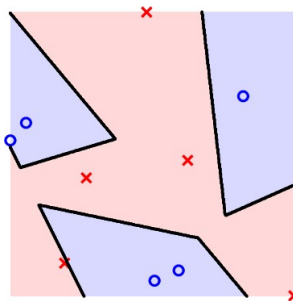
- 1 0
- 2 1
- 3 1126
- 4 5566

Reference Answer: 1


When a feature is a constant, permutation does not change its value. Then, $E_{\text{oob}}(G)$ and $E_{\text{oob}}^{(p)}(G)$ are the same, and thus $\text{importance}(i) = 0$.

A Simple Data Set


$g_{\text{C\&RT}}$
with random combination



$g_t (N' = N/2)$

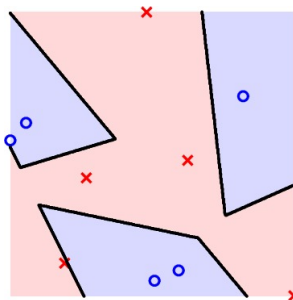


G with first t trees




A Simple Data Set


$g_{\text{C\&RT}}$
with random combination



$g_t (N' = N/2)$

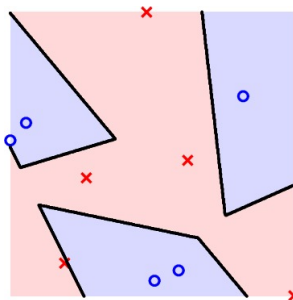


G with first t trees




A Simple Data Set


$g_{\text{C\&RT}}$
with random combination



$g_t (N' = N/2)$

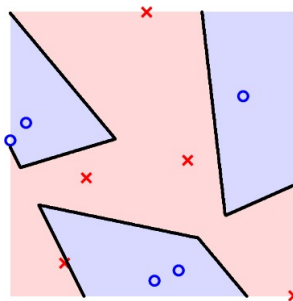


G with first t trees

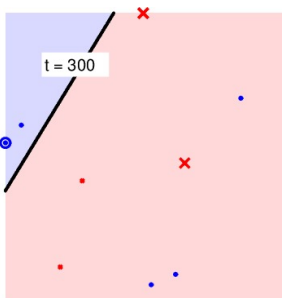


A Simple Data Set

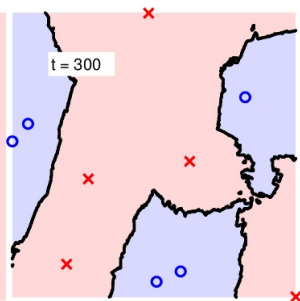
$g_{\text{C\&RT}}$
with random combination



$g_t (N' = N/2)$




G with first t trees




A Simple Data Set


$g_{\text{C\&RT}}$
with random combination



$g_t (N' = N/2)$

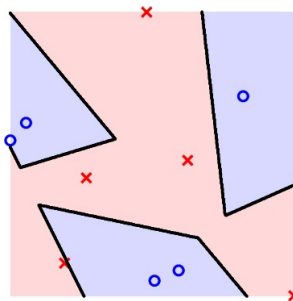


G with first t trees

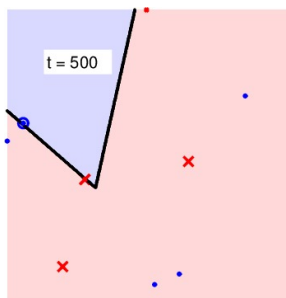


A Simple Data Set


$g_{\text{C\&RT}}$
with random combination



$g_t (N' = N/2)$




G with first t trees




A Simple Data Set


$g_{\text{C\&RT}}$
with random combination



$g_t (N' = N/2)$

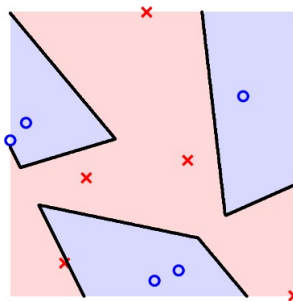


G with first t trees

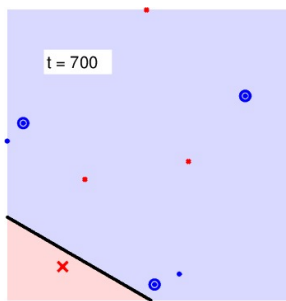


A Simple Data Set

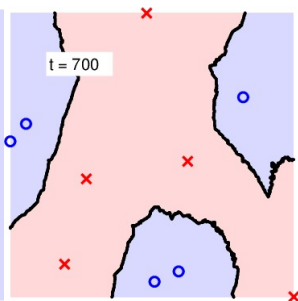
$g_{\text{C\&RT}}$
with random combination



$g_t (N' = N/2)$

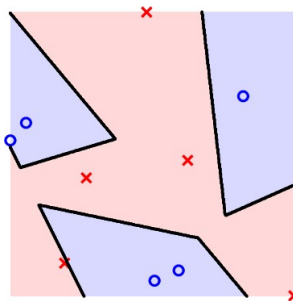


G with first t trees

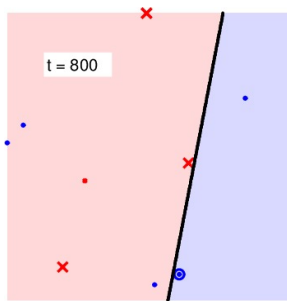


A Simple Data Set

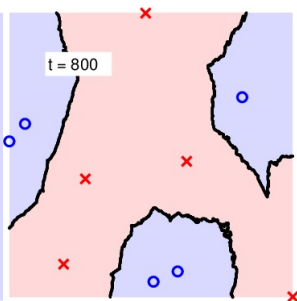
$g_{\text{C\&RT}}$
with random combination



$g_t (N' = N/2)$

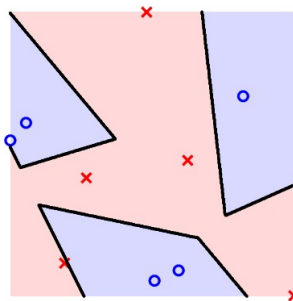


G with first t trees




A Simple Data Set


$g_{\text{C\&RT}}$
with random combination



$g_t (N' = N/2)$

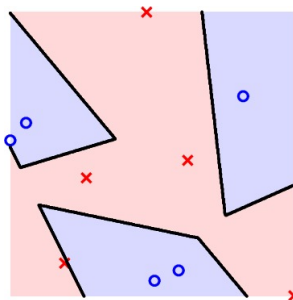


G with first t trees

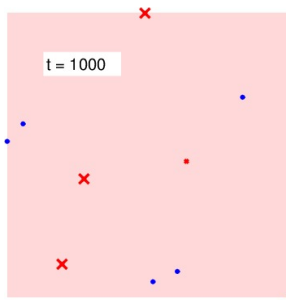


A Simple Data Set


$g_{\text{C\&RT}}$
with random combination



$g_t (N' = N/2)$





G with first t trees





**‘smooth’ and large-margin-like boundary
with many trees**



A Complicated Data Set





A Complicated Data Set





A Complicated Data Set



A Complicated Data Set

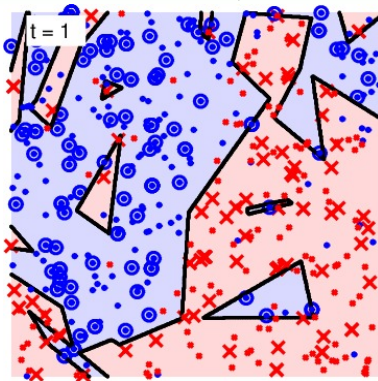
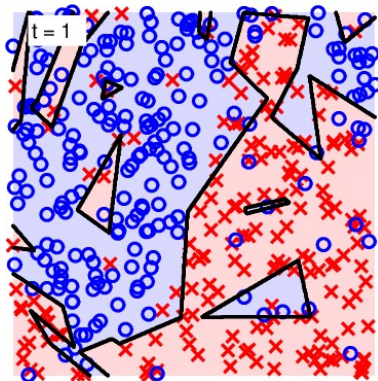


A Complicated Data Set

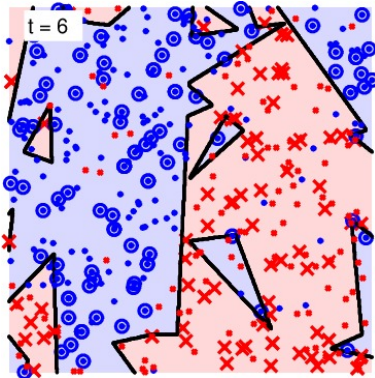
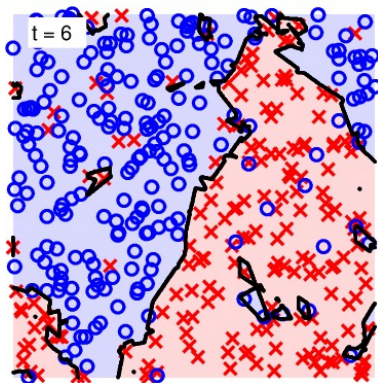


'easy yet robust' nonlinear model

A Complicated and Noisy Data Set

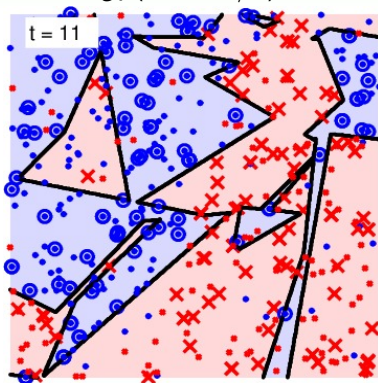
 $g_t (N' = N/2)$  G with first t trees

A Complicated and Noisy Data Set


 $g_t (N' = N/2)$  G with first t trees

A Complicated and Noisy Data Set



$g_t (N' = N/2)$





G with first t trees



A Complicated and Noisy Data Set



A Complicated and Noisy Data Set



noise corrected by voting

How Many Trees Needed?

almost every theory: the more, **the 'better'**
assuming **good** $\bar{g} = \lim_{T \rightarrow \infty} G$

Our NTU Experience

- KDDCup 2013 Track 1 (**yes, NTU is world champion again! :-)**): predicting author-paper relation
- E_{val} of **thousands** of trees: [0.015, 0.019] depending **on seed**;
 E_{out} of top 20 teams: [0.014, 0.019]
- decision: take **12000 trees** with **seed 1**

cons of RF: may need lots of trees **if the whole random process too unstable**
—should double-check **stability of G**
to ensure **enough trees**

Fun Time

Which of the following is **not** the best use of Random Forest?

- ① train each tree with bootstrapped data
- ② use E_{oob} to validate the performance
- ③ conduct feature selection with permutation test
- ④ fix the number of trees, T , to the lucky number 1126

Fun Time

Which of the following is **not** the best use of Random Forest?

- ① train each tree with bootstrapped data
- ② use E_{oob} to validate the performance
- ③ conduct feature selection with permutation test
- ④ fix the number of trees, T , to the lucky number 1126

Reference Answer: ④

A good value of T can depend on the nature of the data and the stability of the whole random process.

Summary

- 1 Embedding Numerous Features: Kernel Models
- 2 Combining Predictive Features: Aggregation Models

Lecture 10: Random Forest

- Random Forest Algorithm
bag of trees on randomly projected subspaces
- Out-Of-Bag Estimate
self-validation with OOB examples
- Feature Selection
permutation test for feature importance
- Random Forest in Action
'smooth' boundary with many trees

- **next: boosted decision trees beyond classification**

- 3 Distilling Implicit Features: Extraction Models